
Homework 4 — Graph algorithms

Sections in CLRS: pp. 525–530 (repeat section B.4 if necessary)

Exercises: 22.1–1

Sections in CLRS: pp. 531–534 (up to "shortest paths")

- We skip the correctness proofs and instead concentrate on breadth-first search and its properties.
- The colors GRAY, WHITE and BLACK are mostly used to argue about the correctness of breadth-first search and are not necessary in an implementation. In addition, we will not build predecessor graphs so you can ignore the π variable in the algorithm. This leads to the following simplified algorithm which uses $d = \infty$ instead of WHITE.

```
BFS( $s$ )
  for each  $u \in V - \{s\}$ 
    do  $d[u] \leftarrow \infty$ 
   $d[s] \leftarrow 0$ 
   $Q \leftarrow \{s\}$ 
  while  $Q \neq \emptyset$ 
    do remove  $u$  from  $Q$ 
    for each  $v \in Adj[u]$ 
      do if  $d[v] = \infty$ 
        then  $d[v] \leftarrow d[u] + 1$ 
        put  $v$  onto  $Q$ 
```

- Note that because one is interested in the distance from the start vertex s , BFS doesn't necessarily visit all the vertices in the graph.


Exercises: 22.2–2, 22.2–3

Sections in CLRS: pp. 540–543, 546–547


- Again, ignore the predecessor graphs (π) and proof of correctness. Follow the algorithm in the book (page 478) but ignore line 3 in DFS and line 6 in DFS-VISIT.
- Note that DFS visits all vertices due to the outer loop. The results can however vary depending on how vertices are selected.
- For a directed graph, edges are marked as belonging to one of the following categories: tree edges, back edges, forward edges and cross edges. For undirected graphs there are only tree and back edges.

Exercises: 22.3–2, 22.3–7

- For 22.3–2 you do not have to produce a sequence of pictures like in Figure 22.4 in the book: it is enough to mark $d[u], f[u]$ for each vertex u and to classify the edges as described above. You do not even have to color the vertices if you write $d[u], f[u]$ next to each vertex. For instance, a gray vertex would have $d[u]$ defined, but not $f[u]$.



WHITE




GRAY




BLACK

The edge category can be summarized with the following situations:

Tree edge




Back edge



$d[u] < d[v]$: Forward edge (F)

$d[u] > d[v]$: Cross edge (C)



Sections in CLRS: pp. 549–551

Exercises: 22.4–1, 22.4–3, 22.4–5

- If you succeed in giving an algorithm for 22.4–3 but think that it runs in $O(V + E)$, chances are your analysis is not tight enough. Consider an acyclic, undirected forest with $|E| \leq |V| - 1$ (theorem B.2, page 1085 in the book). How many edges will your algorithm visit in the worst case?

Sections in CLRS: We will not consider section 22.5 “Strongly connected components” but you should know what a SCC is (see also page 1082) and that they can be identified in $O(V + E)$ time.

Generally speaking, breadth and depth-first search are useful algorithms because they produce useful information for other algorithms like topological sorting, identifying planar graphs (those that can be drawn on a paper without intersecting edges), etc.




Figure 1: An undirected graph.

Sections in CLRS: pp. 561–566

Exercises: 23.1–1

Sections in CLRS: 570–573 (Prim’s algorithm)

- If MST-PRIM is difficult to understand you can study the following “higher-level” version of Prim’s algorithm first.

MST-PRIM(G, w, r)

▷ Suppose that $V = \{1, 2, \dots, n\}$.

▷ $(U, V - U)$ partitions the graph in two halves.

$U \leftarrow \{1\}$

while $U \neq V$

do let (u, v) be the cheapest edge such that $u \in U$ and $v \in (V - U)$.

$\pi[v] \leftarrow u$ ▷ u is a parent of v in the MST

$U \leftarrow U \cup \{v\}$

return π ▷ return the MST

The trick is of course to repeatedly find the cheapest edge crossing $(U, V - U)$ while the border is moving along. The book employs an efficient way to find the cheapest edge by having a priority Q , sorted by the cost of those edges crossing $(U, V - U)$.

Exercises: Find a minimum spanning tree for the graph in figure 1 using Prim’s algorithm.

Sections in CLRS: Read pp. 580–583, skip “representing shortest paths” and continue with pp. 585–587. Then study Dijkstra’s algorithm, pp. 595–599, but don’t go through the proof until you have done exercise 24.3–1.

Exercises: 24.3–1, 24.3–2, 24.3–3.

Sections in CLRS: Read pp. 620–627. Again you do not need to concern yourself with representing shortest paths (page 621, line 11 and down to the bottom).

Exercises: 25.1–2. Think about why FASTER-ALL-PAIRS-SHORTEST-PATHS works even if n is not an exact power-of-two.

Sections in CLRS: Study pp. 629–632 (up to “Constructing a shortest path”) carefully.

Exercises: 25.2–1, 25.2–4, 25.2–6.

- Exercise 25.2–1 may seem like a lot of work, but will give you the proper insight into solving the other two exercises.

Also show the cheapest path from vertex 3 to vertex 1.

- Exercise 25.2–4 describes the version of Floyd-Warshall’s algorithm that is used in practice. The important question here is whether you can do it using only *one* matrix D .
- If you get stuck on 25.2–6 I would suggest that you construct a small example with a negative cycle.